

CS3211 Tutorial 5

Goroutines and Channels
Simon

Adapted From Sriram's Slides

Why Go?

Message Passing as a First-Class Citizen


- Shared memory \Rightarrow problem

- So: no shared memory \Rightarrow no problem?


- Hah. (also $P \Rightarrow Q \nleftrightarrow !P \Rightarrow !Q$)

- Why do we use channel when we have mutexes?

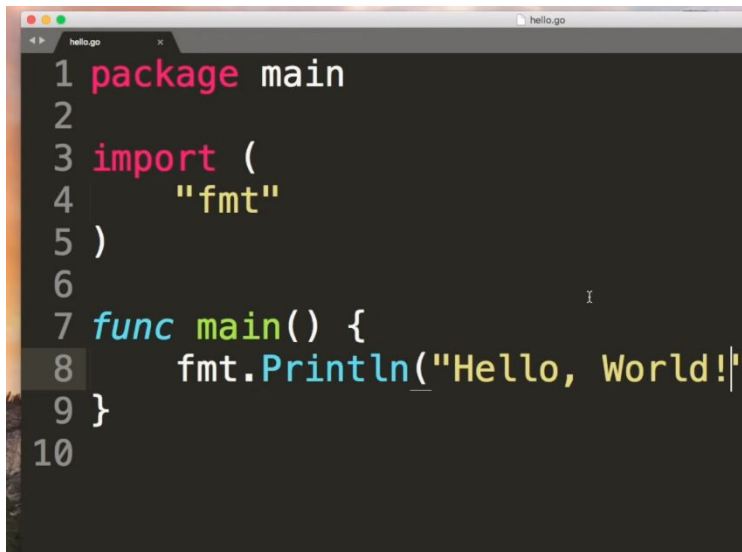
- Because it provides a more elegant concurrency structure (argued by CSP)
 - *Do not communicate by sharing memory; instead, share memory by communicating. (Effective Go)*
 - <https://go.dev/blog/codelab-share>
 - <http://www.usingcsp.com/>



Golang

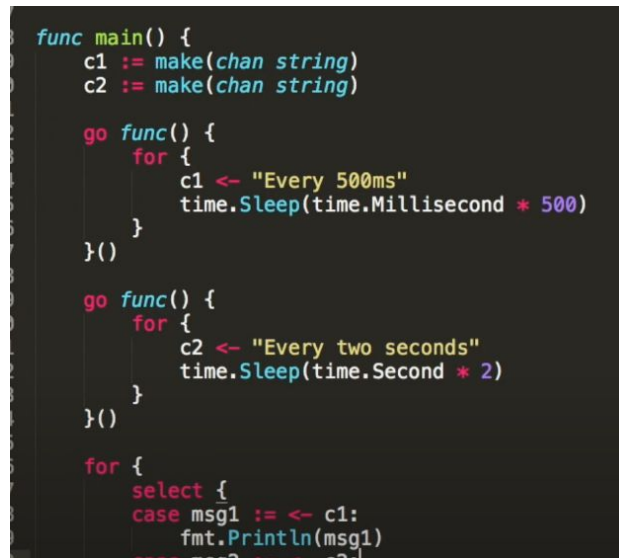


Two Useful Videos



```
1 package main
2
3 import (
4     "fmt"
5 )
6
7 func main() {
8     fmt.Println("Hello, World!")
9 }
10
```

[Learn Go in 12 Minutes](#)



```
func main() {
    c1 := make(chan string)
    c2 := make(chan string)

    go func() {
        for {
            c1 <- "Every 500ms"
            time.Sleep(time.Millisecond * 500)
        }
    }()

    go func() {
        for {
            c2 <- "Every two seconds"
            time.Sleep(time.Second * 2)
        }
    }()

    for {
        select {
            case msg1 := <- c1:
                fmt.Println(msg1)
            case msg2 := <- c2:
                fmt.Println(msg2)
        }
    }
}
```


[Concurrency in Go](#)

Go's Special Sauce: Inbuilt...

- Channels
- Supporting characters (non-exhaustive)
 - Waitgroups
 - Range, (v, ok)...
 - select, select-default
 - Defer

Go Channels

- Essentially, a native **thread-safe** multi-producer, multi-consumer “queue”!



Exploring Channels

- What is the output of this program? [p]

```
1  package main
2
3  import "fmt"
4
5  func main() {
6      c := make(chan int)
7      c <- 1
8      fmt.Println(<-c)
9  }
10 |
```

Exploring Channels

- What is the output of this program? [p]
- Deadlock!
 - Main() blocks on channel send since no one to receive yet
 - Any simple ways to avoid this?

<https://fsmbolt.comp.nus.edu.sg/z/nMP766>

```
1  package main
2
3  import "fmt"
4
5  func main() {
6      c := make(chan int)
7      c <- 1
8      fmt.Println(<-c)
9  }
10 |
```


Exploring Channels

- How many possible outputs does this program have, if any? [p]

```
1  package main
2
3  import (
4      |   "fmt"
5  )
6
7  func producer(val int, c chan<- int) {
8      |   c <- val
9      |   c <- val
10 }
11
12
13 func main() {
14     |   c := make(chan int)
15
16     |   go producer(1, c)
17     |   go producer(7, c)
18
19     |   fmt.Println(<-c + <-c)
20 }
```

Exploring Channels

- How many possible outputs does this program have, if any? [p]
- 3 possible outputs
 - 1 + 1
 - 1 + 7 (or 7 + 1)
 - 7 + 7

<https://fsmbolt.comp.nus.edu.sg/z/rhfacz>

```
1  package main
2
3  import (
4      "fmt"
5  )
6
7  func producer(val int, c chan<- int) {
8      c <- val
9      c <- val
10 }
11
12
13 func main() {
14     c := make(chan int)
15
16     go producer(1, c)
17     go producer(7, c)
18
19     fmt.Println(<-c + <-c)
20 }
```

Why does this matter?

- Channels are not an automatic solution to problems
- They can also be complex in their own ways!

Tutorial Q1.1 – 1.2

1.1 Concurrent Counter: Waitgroups

- Task: multiple threads adding to a counter
- Initial answer: <https://fsmbolt.comp.nus.edu.sg/z/ch951W>
- **Problem: We're not waiting for all goroutines to finish!**

1.1 Concurrent Counter

<https://fsmbolt.comp.nus.edu.sg/z/r347WE>

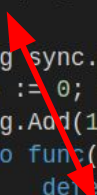
- Concept of `sync.Waitgroup`
- A way for a thread to wait for some number of threads to finish (but more general than this)
- **Problem?**

```
1 package main
2
3 import (
4     "fmt"
5     "sync"
6 )
7
8 func main() {
9     count := 0
10
11     var wg sync.WaitGroup
12     for i := 0; i < 1000; i++ {
13         wg.Add(1) // add BEFORE spawning
14         go func() {
15             defer wg.Done() // RAII. called on exit
16             count++
17         }()
18     }
19     wg.Wait() // wait until all 1000 goroutines are done
20
21     fmt.Println("Count: ", count)
22 }
23
```

1.1 Concurrent Counter

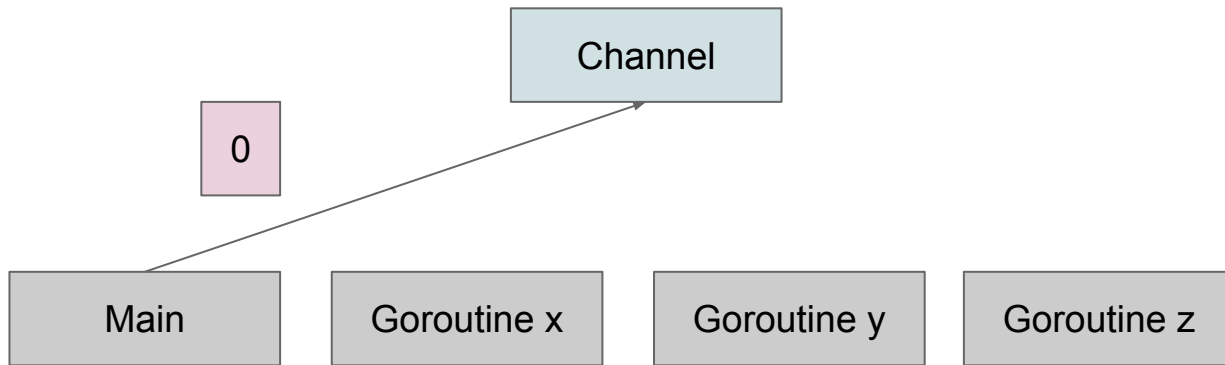
- Concept of `sync.Waitgroup`
- A way for a thread to wait for some number of threads to finish (but more general than this)
- **Problem? Straightforward data race: `count++`-ed by reference across threads**

```
1 package main
2
3 import (
4     "fmt"
5     "sync"
6 )
7
8 func main() {
9     count := 0
10
11     var wg sync.WaitGroup
12     for i := 0; i < 1000; i++ {
13         wg.Add(1) // add BEFORE spawning
14         go func() {
15             defer wg.Done() // RAII. called on exit
16             count++
17         }()
18     }
19     wg.Wait() // wait until all 1000 goroutines are done
20
21     fmt.Println("Count: ", count)
22 }
23
```



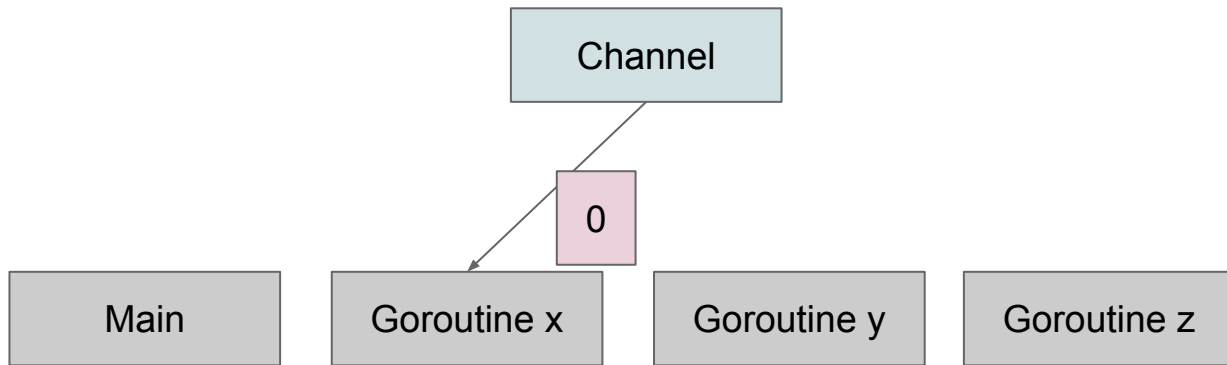
1.1 Concurrent Counter: Solution 1

- “Channel solution”: have an **exclusive owner** of count each time



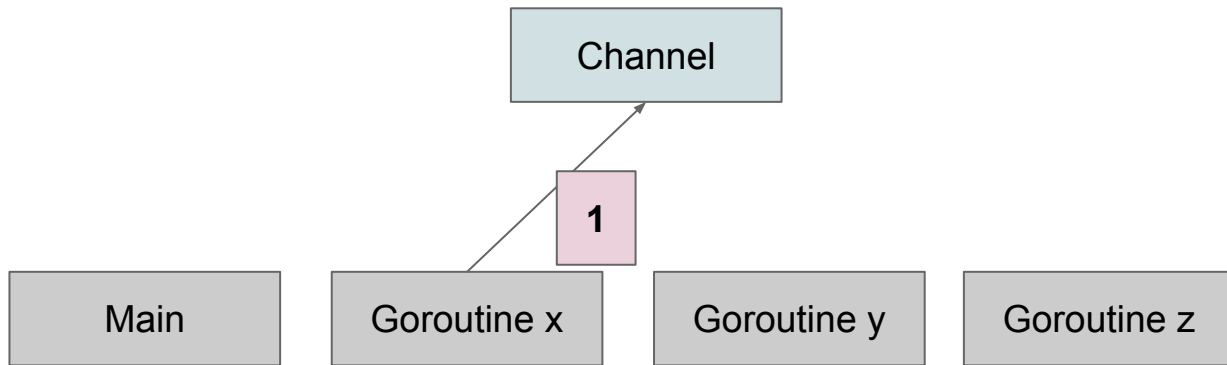
1.1 Concurrent Counter: Solution 1

- “Channel solution”: have an **exclusive owner** of count each time



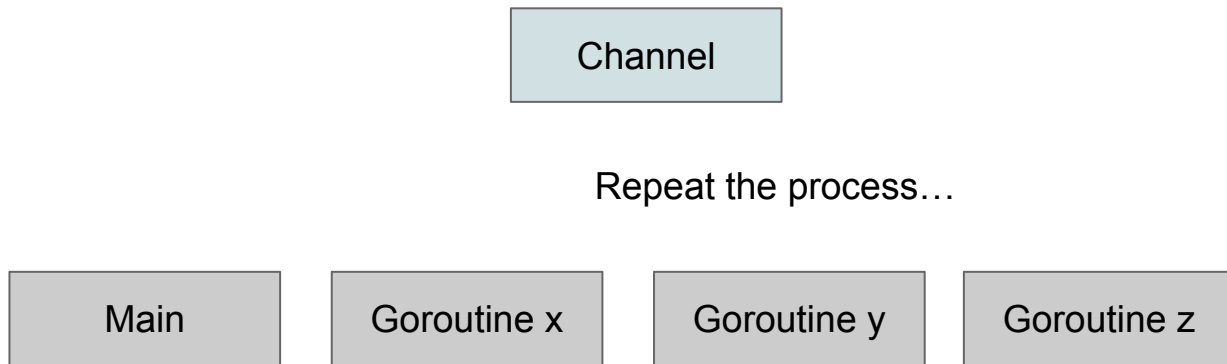
1.1 Concurrent Counter: Solution 1

- “Channel solution”: have an **exclusive owner** of count each time



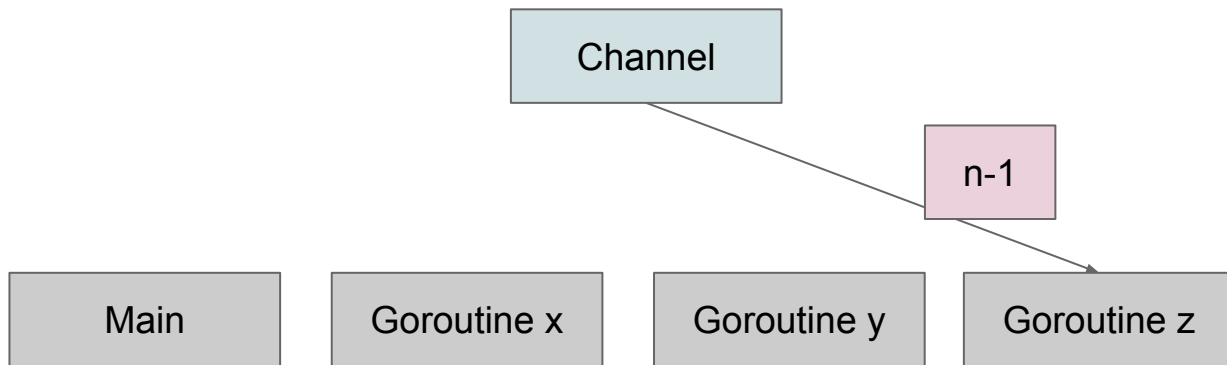
1.1 Concurrent Counter: Solution 1

- “Channel solution”: have an **exclusive owner** of count each time



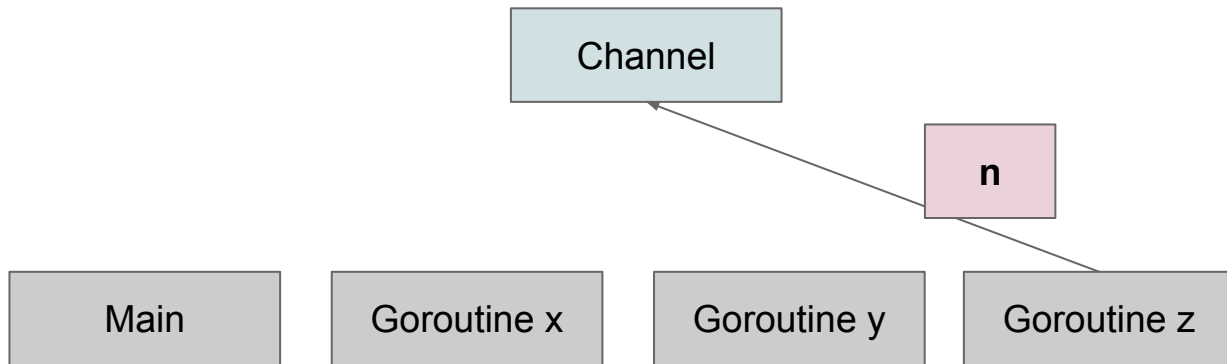
1.1 Concurrent Counter: Solution 1

- “Channel solution”: have an **exclusive owner** of count each time



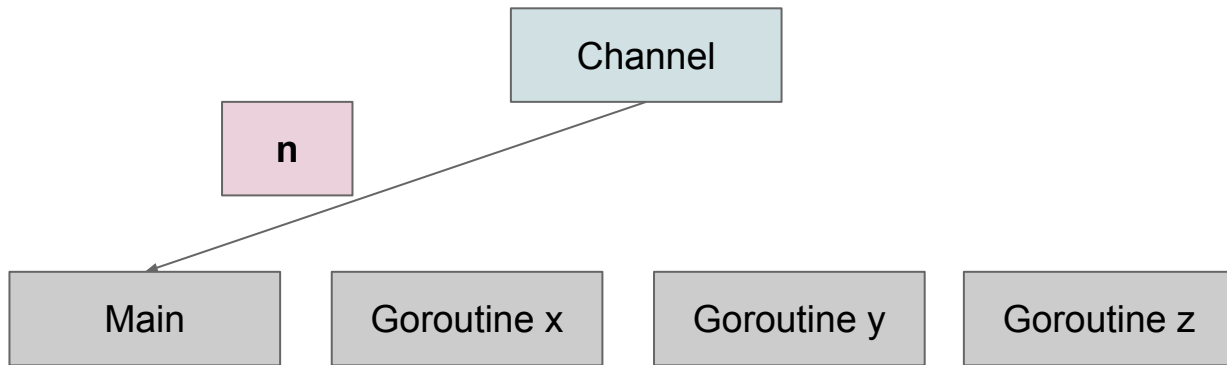
1.1 Concurrent Counter: Solution 1

- “Channel solution”: have an **exclusive owner** of count each time



1.1 Concurrent Counter: Solution 1

- “Channel solution”: have an **exclusive owner** of count each time



Also, any non-correctness-related issues with this situation?

<https://fsmbolt.comp.nus.edu.sg/z/ao13f7>

Which line is stuck (deadlock)?

```
1  package main
2
3  import (
4      "fmt"
5      "sync"
6  )
7
8  func main() {
9      ch := make(chan int) // make "unbuffered" channel
10     var wg sync.WaitGroup
11     for i := 0; i < 1000; i++ {
12         wg.Add(1)
13         go func() {
14             defer wg.Done()
15             count := <-ch // blocking dequeue
16             count++      // safely add 1 as the exclusive owner
17             ch <- count   // blocking enqueue (for another consumer)
18         }()
19     }
20     ch <- 0 // main sends initial value; blocking enqueue
21
22     wg.Wait() // wait for all goroutines
23     fmt.Println("Count: ", <-ch) // dequeue final result
24 }
```

Which line is stuck (deadlock)?

```
1  package main
2
3  import (
4      "fmt"
5      "sync"
6  )
7
8  func main() {
9      ch := make(chan int) // make "unbuffered" channel
10     var wg sync.WaitGroup
11     for i := 0; i < 1000; i++ {
12         wg.Add(1)
13         go func() {
14             defer wg.Done()
15             count := <-ch // blocking dequeue
16             count++      // safely add 1 as the exclusive owner
17             ch <- count   // blocking enqueue (for another consumer)
18         }()
19     }
20     ch <- 0 // main sends initial value; blocking enqueue
21
22     wg.Wait() // wait for all goroutines
23     fmt.Println("Count: ", <-ch) // dequeue final result
24 }
```

1000th goroutine tries to
write 1000 to channel

Which line is stuck (deadlock)?

```
1  package main
2
3  import (
4      "fmt"
5      "sync"
6  )
7
8  func main() {
9      ch := make(chan int) // make "unbuffered" channel
10     var wg sync.WaitGroup
11     for i := 0; i < 1000; i++ {
12         wg.Add(1)
13         go func() {
14             defer wg.Done()
15             count := <-ch // blocking dequeue
16             count++      // safely add 1 as the exclusive owner
17             ch <- count   // blocking enqueue (for another consumer)
18         }()
19     }
20     ch <- 0 // main sends initial value; blocking enqueue
21
22     wg.Wait() // wait for all goroutines
23     fmt.Println("Count: ", <-ch) // dequeue final result
24 }
```

1000th goroutine tries to
write 1000 to channel

But main is stuck here!
Not receiving

Which line is stuck (deadlock)?

Defer only runs after the go func() is done here!

1000th goroutine tries to write 1000 to channel

But main is stuck here!
Not receiving

```
1  package main
2
3  import (
4      "fmt"
5      "sync"
6  )
7
8  func main() {
9      ch := make(chan int) // make "unbuffered" channel
10     var wg sync.WaitGroup
11     for i := 0; i < 1000; i++ {
12         wg.Add(1)
13         go func() {
14             defer wg.Done()
15             count := <-ch // blocking dequeue
16             count++      // safely add 1 as the exclusive owner
17             ch <- count  // blocking enqueue (for another consumer)
18         }()
19     }
20     ch <- 0 // main sends initial value; blocking enqueue
21
22     wg.Wait() // wait for all goroutines
23     fmt.Println("Count: ", <-ch) // dequeue final result
24 }
```

1.2 Solution 1: **wg.Done()** before channel send

- “Natural” use of *defer* really messed us up!
- No easy day in concurrency land...

```
8 func main() {
9     ch := make(chan int) // make unbuffered channel
10    var wg sync.WaitGroup
11    for i := 0; i < 1000; i++ {
12        wg.Add(1)
13        go func() {
14            count := <-ch
15            count++
16            wg.Done() // B
17            ch <- count // A
18        }()
19    }
20    // main sends initial value; block until received
21    ch <- 0
22    wg.Wait() // C
23    fmt.Println("Count: ", <-ch) // D
24 }
```

1.2 Solution 2: buffered channel

- Scary to generalize this...

```
8 func main() {
9     ...ch := make(chan int, 1) // make buffered channel of size 1
10    var wg sync.WaitGroup
11    for i := 0; i < 1000; i++ {
12        wg.Add(1)
13        go func() {
14            defer wg.Done() // defer to release
15            // read count, increment, send away
16            count := <-ch // blocked until sent
17            count++
18            ch <- count
19        }()
20    }
21    ch <- 0 // main sends initial value; block until received
22    wg.Wait()
23    fmt.Println("Count: ", <-ch)
24 }
```

1.2 Solution 2: buffered channel

- What happens here? [p]

```
8 func main() {
9     ...ch := make(chan int, 2) // make buffered channel of size 2
10    var wg sync.WaitGroup
11    for i := 0; i < 1000; i++ {
12        wg.Add(1)
13        go func() {
14            defer wg.Done() // defer to release
15            // read count, increment, send away
16            count := <-ch // blocked until sent
17            count++
18            ch <- count
19        }()
20    }
21    ch <- 0 // main sends initial value; block until received
22    wg.Wait()
23    fmt.Println("Count: ", <-ch)
24 }
```

1.2 Solution 2: buffered channel

- What happens here? [p]
- Thankfully, still correct (outputs 1000). But like.. why?

<https://fsmbolt.comp.nus.edu.sg/z/cTWvz5>

```
8 func main() {
9     ch := make(chan int, 2) // make buffered channel of size 2
10    var wg sync.WaitGroup
11    for i := 0; i < 1000; i++ {
12        wg.Add(1)
13        go func() {
14            defer wg.Done() // defer to release
15            // read count, increment, send away
16            count := <-ch // blocked until sent
17            count++
18            ch <- count
19        }()
20    }
21    ch <- 0 // main sends initial value; block until received
22    wg.Wait()
23    fmt.Println("Count: ", <-ch)
24 }
```

1.2 Solution 2: buffered channel

- What happens here? [p] <https://fsmbolt.comp.nus.edu.sg/z/cTWvz5>
- Thankfully, still correct. Max 1 value in the queue based on our design, but not restricted by the channel. Had to reason about it...

```
8 func main() {
9     ch := make(chan int, 2) // make buffered channel of size 2
10    var wg sync.WaitGroup
11    for i := 0; i < 1000; i++ {
12        wg.Add(1)
13        go func() {
14            defer wg.Done() // defer to release
15            // read count, increment, send away
16            count := <-ch // blocked until sent
17            count++
18            ch <- count
19        }()
20    }
21    ch <- 0 // main sends initial value; block until received
22    wg.Wait()
23    fmt.Println("Count: ", <-ch)
24 }
```

Tutorial Q1.3

1.3 *Independently* producing / merging counts

- We don't want each thread to take *exclusive* access of a global count (this is basically a mutex implemented by channels!)
- How can we parallelize adding to the final count?

1.3 *Independently* producing / merging counts

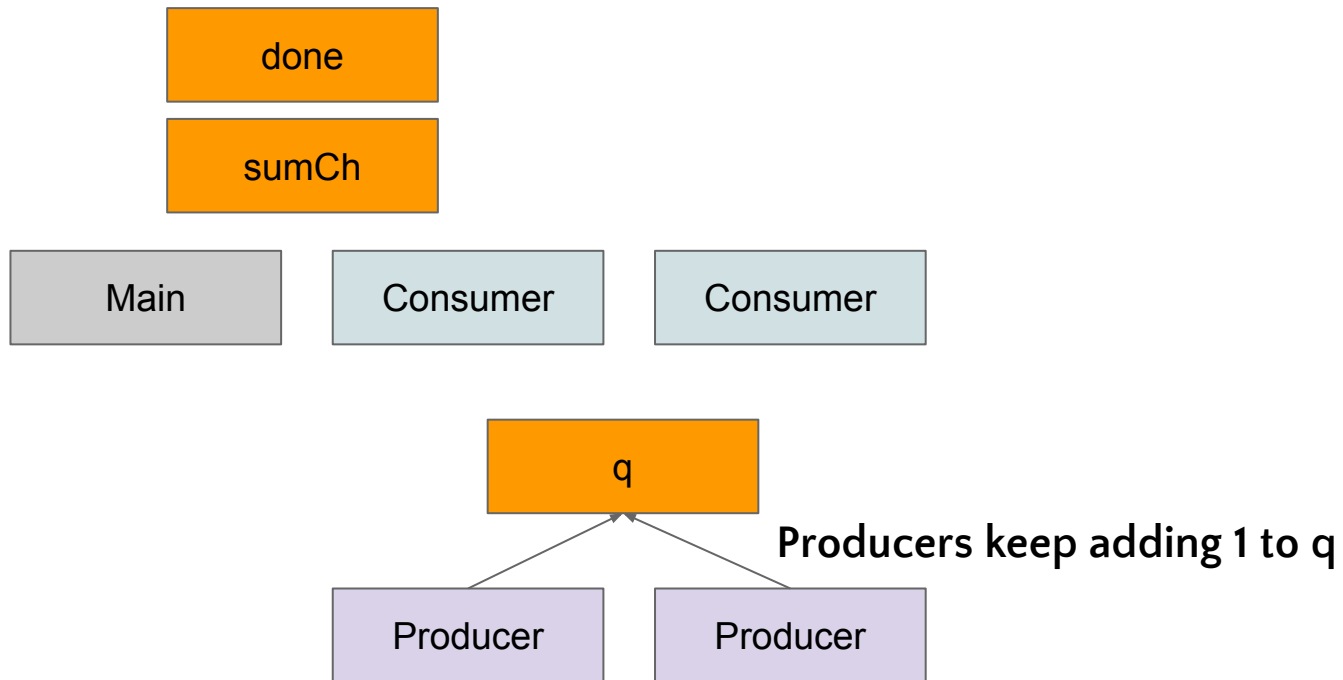
- We don't want each thread to take *exclusive* access of a global count (this is basically a mutex implemented by channels!)
- How can we parallelize adding to the final count?
- Step 1: Consumers read from producers and update local counts
<https://fsmbolt.comp.nus.edu.sg/z/7Yh9sc>

1.3 *Independently* producing / merging counts

- We don't want each thread to take *exclusive* access of a global count (this is basically a mutex implemented by channels!)
- How can we parallelize adding to the final count?
- Step 1: Consumers read from producers and update local counts
<https://fsmbolt.comp.nus.edu.sg/z/7Yh9sc>
- Step 2: Consumers send local counts to overall sum channel when production is done
<https://fsmbolt.comp.nus.edu.sg/z/nPvfsq>

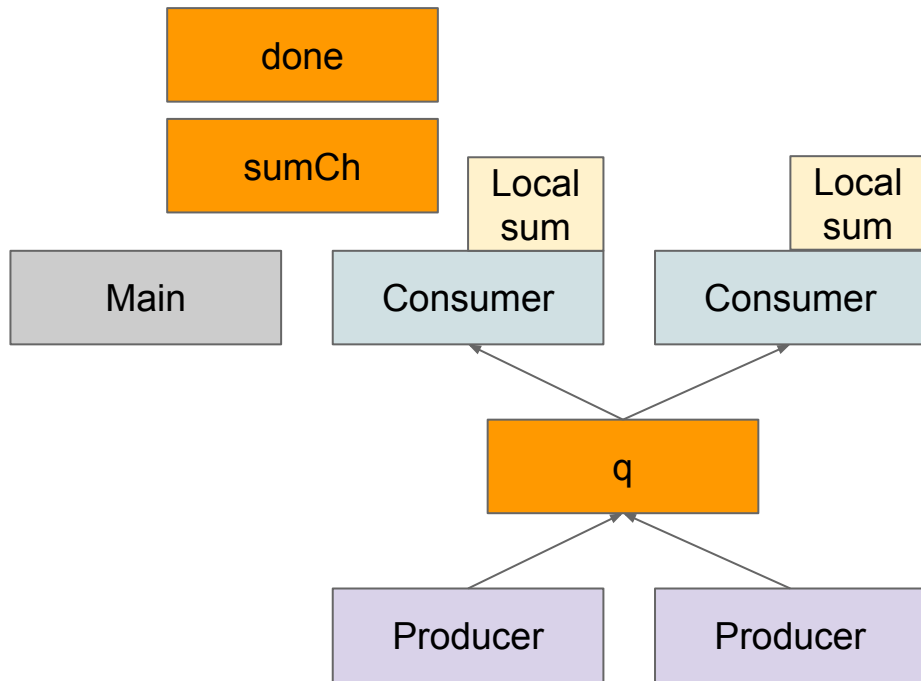
1.3 *Independently* producing / merging counts

Overall idea:



1.3 *Independently* producing / merging counts

Overall idea:

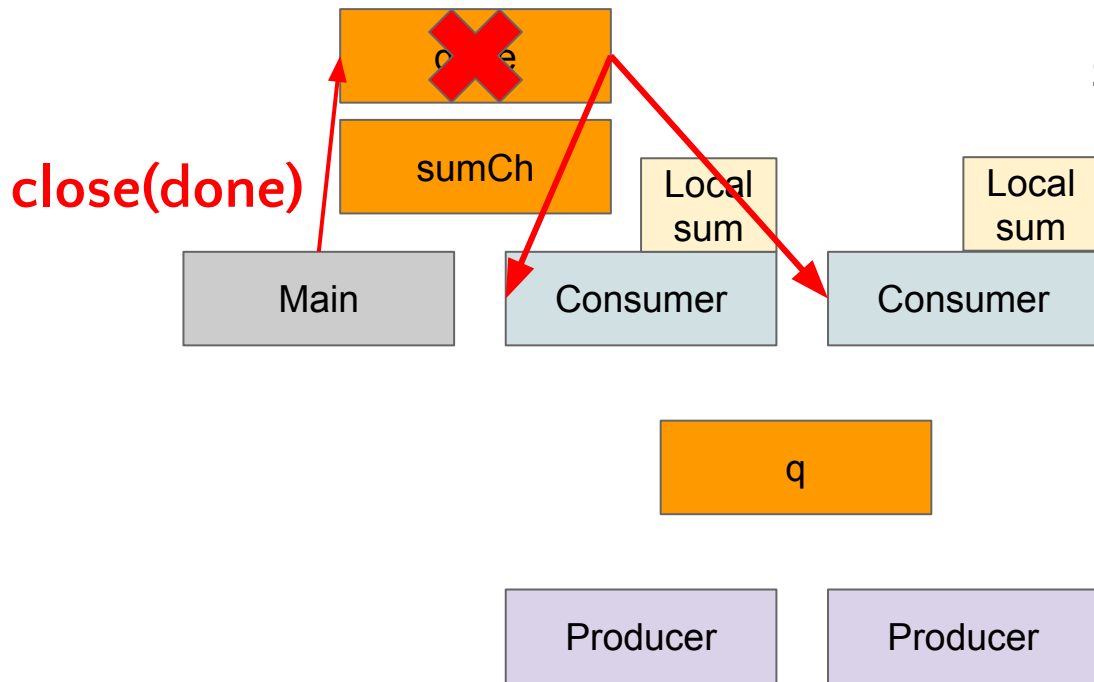


Consumers take 1s off q and
add to local count

1.3 *Independently* producing / merging counts

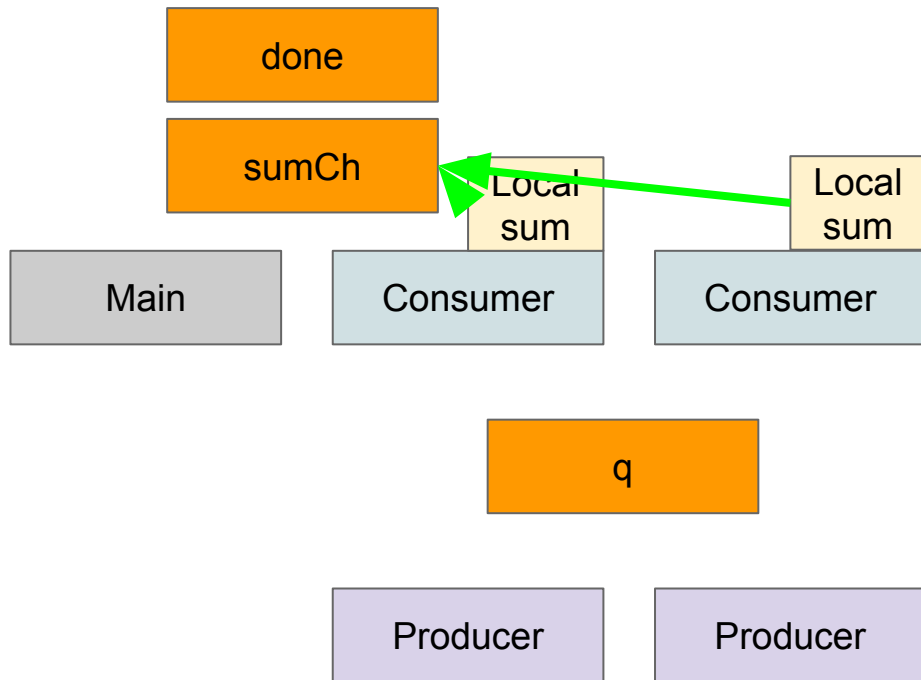
Overall idea:

When consumers see that
done channel is closed...
(closed channel will return
zero-value of type when read)



1.3 *Independently* producing / merging counts

Overall idea:



They send their local sum
to the sumCh channel

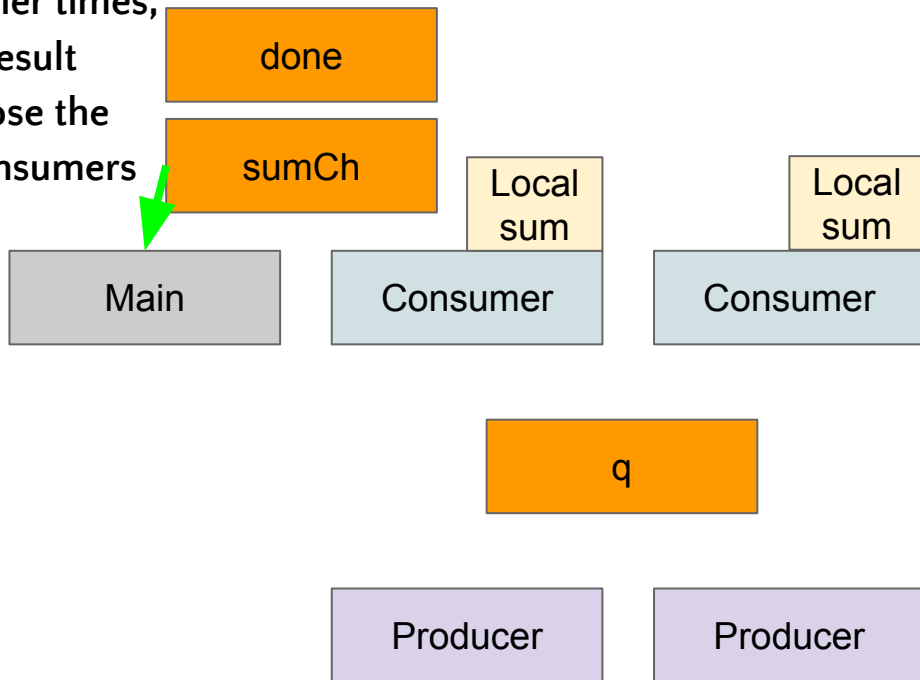
1.3 *Independently* producing / merging counts

Overall idea:

Read NumConsumer times,

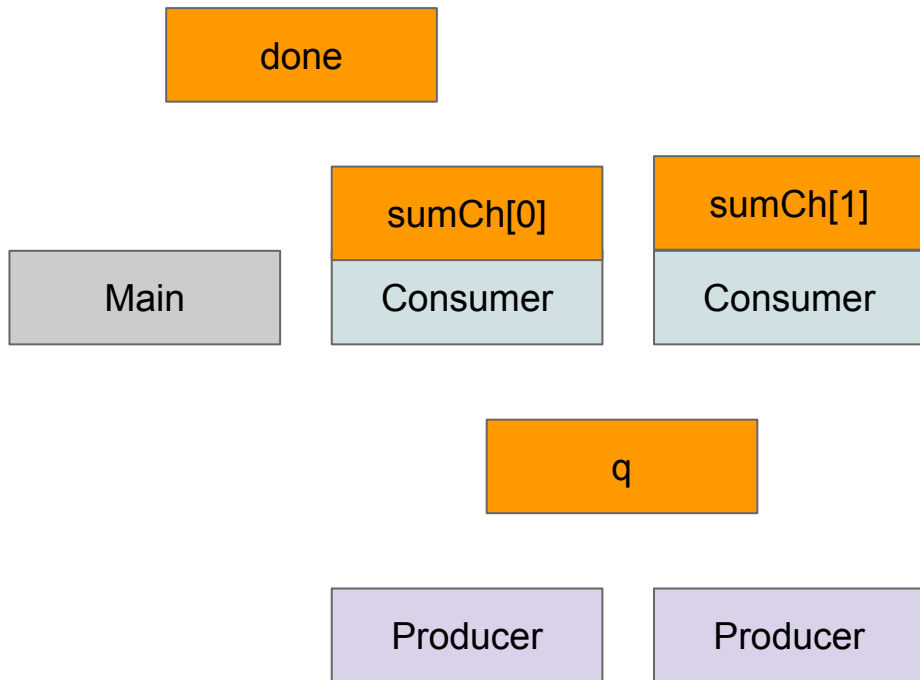
Add to final result

Qn: can you close the
channel from consumers



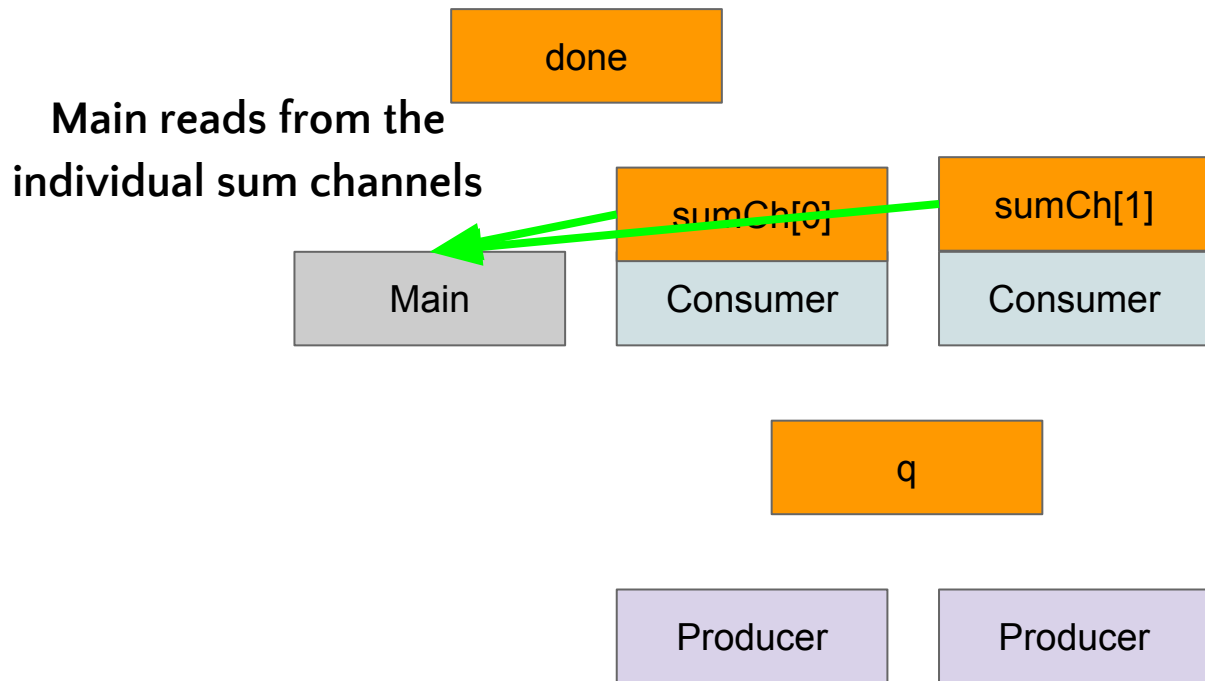
1.3 *Independently* producing / merging counts

Alternate design: separate sum channels for each consumer!



1.3 *Independently* producing / merging counts

Alternate design: separate sum channels for each consumer!



Tutorial Q2

Quick Detour to Go Patterns

Extra Qn: Is it possible to make Lock Free Queue using Go Channels?

For-Select Go Pattern

```
for {  
    select {  
        case q <- 1: // keeps sending to q  
        case <-done: // should exit  
            return  
    }  
}
```

Note that Goroutines are resources too, failing to exit the Goroutines may cause memory leaks

Select Default Go Pattern

```
select {  
    case q <- num:  
        return true  
    default:  
}  
return false  
}
```

Qn:

1. What's the difference between C++ switch case and Go select?
2. If you know Linux Select, what's the difference with Go Select?

Context in Go Pattern

The `context` package is used for carrying deadlines, cancellations, and other request-scoped values across goroutines.

```
// Create common context in main goroutine
ctx, _ := context.WithTimeout(context.Background(), time.Second)
...

// Use the context in some Goroutines
for {
    select {
    case q <- 1: ...
    case <-ctx.Done():
        return
    }
}
```

Why does this all matter?

- Channel/message-passing-based thinking is extremely useful for certain kinds of problems, it provides elegant solutions to many concurrency problems
- Avoids many issues with shared memory
- **However: please look at the safety of the channels themselves, there are certain requirements that you need to adhere when you are using channels such as never close a channel twice, etc.**

Extra: Exploring Channels

- What is the final value of this program? (Assume that it correctly compiles) [p]

```
7 // Increment value in channel
8 func consumer(vals <-chan *int, results chan<- int) {
9     val := <-vals
10    *val++
11 }
12
13
14 func main() {
15     // Number of goroutines
16     const num_threads = 700_000
17
18     // Initialize channels and variables
19     v := 0
20     vals := make(chan *int, num_threads)
21     results := make(chan int)
22
23     // Create all consumers
24     for i := 0; i < num_threads; i++ {
25         go consumer(vals, results)
26     }
27
28     // Fill channel with values
29     for i := 0; i < num_threads; i++ {
30         vals <- &v
31     }
32     // ....
33     // ASSUME we wait correctly for all consumers to finish
34     fmt.Println("Final value: ", v)
35 }
```


Extra: Exploring Channels

<https://fsmolt.comp.nus.edu.sg/z/xT1nEz>

- What is the final value of this program? (Assume that it correctly compiles) [p]
- Unknown, much less than 700_000 is very likely
 - Data race!
 - Passed v by *reference* (ptr)
- Channels do not automatically free us from shared memory concerns!

```
7 // Increment value in channel
8 func consumer(vals <-chan *int, results chan<- int) {
9     val := <-vals
10    *val++
11 }
12
13
14 func main() {
15     // Number of goroutines
16     const num_threads = 700_000
17
18     // Initialize channels and variables
19     v := 0
20     vals := make(chan *int, num_threads)
21     results := make(chan int)
22
23     // Create all consumers
24     for i := 0; i < num_threads; i++ {
25         go consumer(vals, results)
26     }
27
28     // Fill channel with values
29     for i := 0; i < num_threads; i++ {
30         vals <- &v
31     }
32     // ....
33     // ASSUME we wait correctly for all consumers to finish
34     fmt.Println("Final value: ", v)
35 }
```

See you next week!