

CS3211 Tutorial 7

Classic Concurrency Problems in C++ & Go

Simon

Adapted from Sriram's and Xue Yong's Slides

H2O Problem

i.e. resource management

1. H2O Problem


- Hydrogen and oxygen atoms arrive at the factory.
- We only want to bond 2 H with 1 O.
- Atoms only start bonding when 2 H and 1 O arrives.
- Other atoms block.
- Only 1 bonding can occur at a time.

Models distributed resources allocation

Sister's Problems: Cigarette Smokers Problem,
River Crossing Problem.

Source:

<https://greenteapress.com/semaphores/LittleBookOfSemaphores.pdf>




H2O Problem: **Shared Memory**

i.e. resource management

1. H₂O Problem

Incoming atoms




factory

completed

1. H₂O Problem

Incoming atoms




factory

completed

1. H₂O Problem

Incoming atoms




factory

completed

1. H₂O Problem


Incoming atoms



factory


completed

1. H₂O Problem




1. H2O Problem

Incoming atoms



factory




completed


Must bond before another bonds!

1. H₂O Problem

Incoming atoms



factory



Solution 1: Using a Barrier

Any problem here?

```
struct WaterFactory {  
    std::barrier<> barrier;  
    WaterFactory() : barrier{3} {}  
  
    void oxygen(void (*bond)()) {  
        barrier.arrive_and_wait();  
        bond();  
    }  
  
    void hydrogen(void (*bond)()) {  
        barrier.arrive_and_wait();  
        bond();  
    }  
};
```

Solution 1: Using a Barrier

```
auto oxygen_bond = []() {
    print("Bonding oxygen\n");
    std::this_thread::sleep_for(std::chrono::milliseconds{5});
    print("Done\n");
};

std::thread t1{[&]() { factory.oxygen(oxygen_bond); }};
std::thread t2{[&]() { factory.oxygen(oxygen_bond); }};
std::thread t3{[&]() { factory.oxygen(oxygen_bond); }};
// Bonding oxygen
// Bonding oxygen
// Bonding oxygen
// Done
// Done
// Done
```

Does not ensure correct types! Here we form ozone rather than water

```
struct WaterFactory {
    std::barrier<> barrier;
    WaterFactory() : barrier{3} {}

    void oxygen(void (*bond)()) {
        barrier.arrive_and_wait();
        bond();
    }

    void hydrogen(void (*bond)()) {
        barrier.arrive_and_wait();
        bond();
    }
};
```

Solution 2: Using Semaphores (Extra)

Any problem here?

```
struct WaterFactory {  
    std::counting_semaphore oSem{1}, hSem{2};  
    WaterFactory () {}  
  
    void oxygen(void (*bond) ()) {  
        oSem.acquire();  
        bond();  
        oSem.release();  
    }  
  
    void hydrogen(void (*bond) ()) {  
        hSem.acquire();  
        bond();  
        hSem.release();  
    }  
};
```

Solution 2: Using Semaphores (Extra)

Any problem here? Yes!

Case 1: 1 or 2 H atoms call
bond() -> we get H₂ or H

Case 2: 1 O atom call bond()
-> We get O

```
struct WaterFactory {
    std::counting_semaphore oSem{1}, hSem{2};
    WaterFactory() {}

    void oxygen(void (*bond)()) {
        oSem.acquire();
        bond();
        oSem.release();
    }

    void hydrogen(void (*bond)()) {
        hSem.acquire();
        bond();
        hSem.release();
    }
};
```

Solution 3: Using Barrier + Semaphores

Any problem here?

```
struct WaterFactory {
    std::counting_semaphore oSem{1}, hSem{2};
    std::barrier<> barrier;
    WaterFactory() : barrier{3} {}

    void oxygen(void (*bond)()) {
        oSem.acquire();
        barrier.arrive_and_wait();
        bond();
        oSem.release();
    }

    void hydrogen(void (*bond)()) {
        hSem.acquire();
        barrier.arrive_and_wait();
        bond();
        hSem.release();
    }
};
```


Solution 3: Using Barrier + Semaphores

Any problem here? **NO!**

Case 1: there are less than 3 atoms -> atoms will block on barrier

Case 2: There are 3+ H atoms -> extra H will block on hSem


Case 3: There are 2+ O atoms -> extra O will block on oSem

```
struct WaterFactory {
    std::counting_semaphore oSem{1}, hSem{2};
    std::barrier<> barrier;
    WaterFactory() : barrier{3} {}


    void oxygen(void (*bond)()) {
        oSem.acquire();
        barrier.arrive_and_wait();
        bond();
        oSem.release();
    }

    void hydrogen(void (*bond)()) {
        hSem.acquire();
        barrier.arrive_and_wait();
        bond();
        hSem.release();
    }
};
```

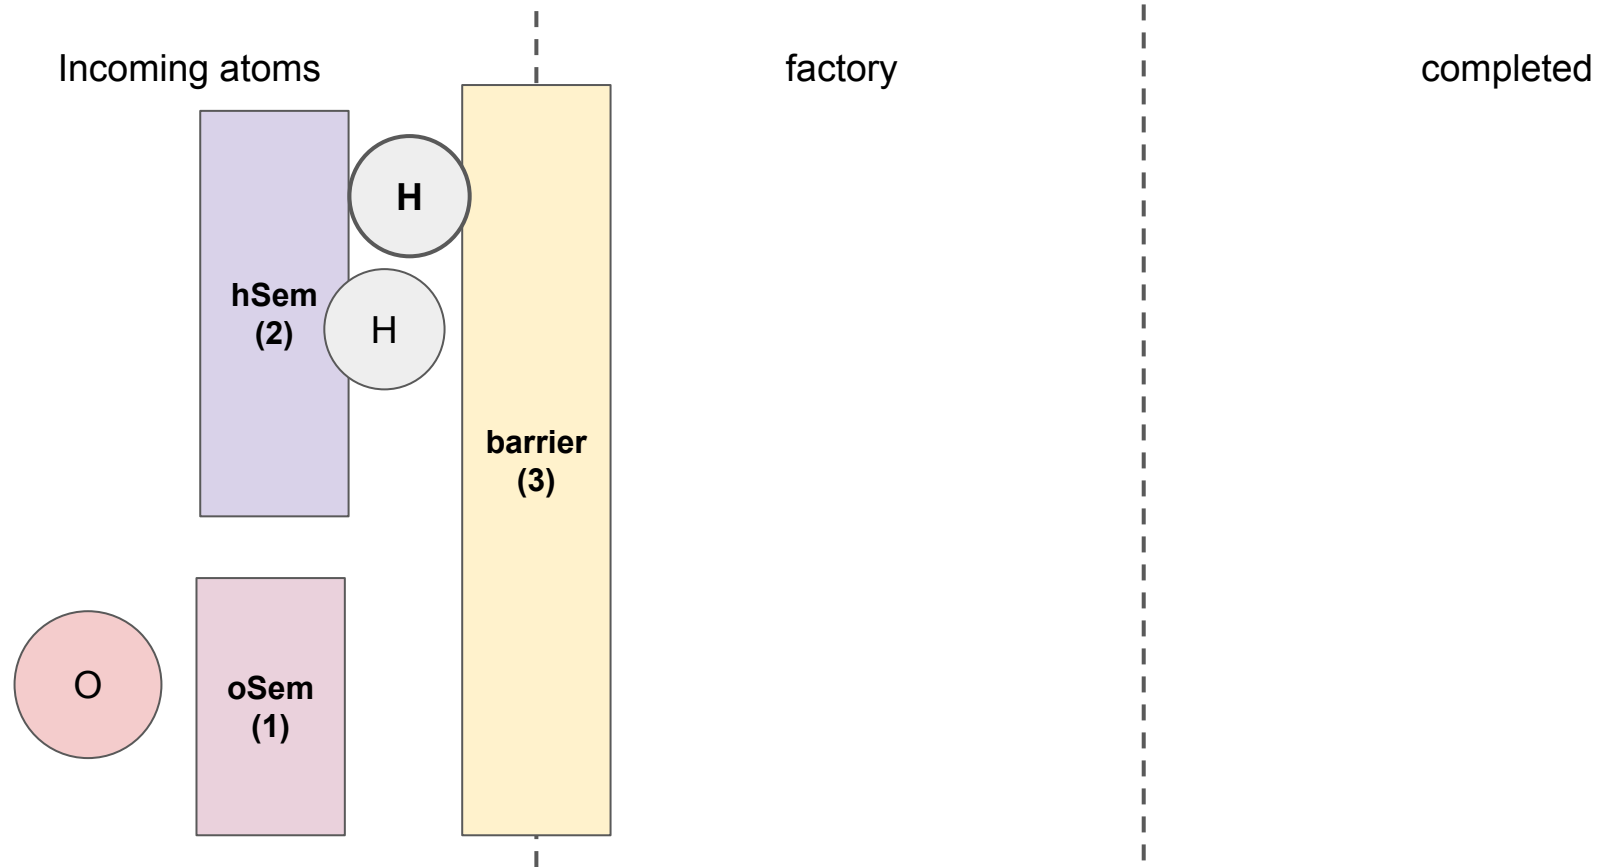
Solution 3: Using Barrier + Semaphores




Solution 3: Using Barrier + Semaphores




Solution 3: Using Barrier + Semaphores



Solution 3: Using Barrier + Semaphores



Solution 3: Using Barrier + Semaphores



Trivia

What if we use an `std::latch` instead of an `std::barrier`?

Trivia


What if we use an `std::latch` instead of an `std::barrier`?

The latch class is a downward counter of type `std::ptrdiff_t` which can be used to synchronize threads. The value of the counter is initialized on creation. Threads may block on the latch until the counter is decremented to zero. There is no possibility to increase or reset the counter, which makes the latch a single-use barrier.

H2O Problem: **Channels**

i.e. resource management

Go Solution 1:



Go Solution 1:

```
// Step 1: (Precommit)
//      Receive arrival requests from 2 hydrogen and 1 oxygen atoms
h1 := <-wfd.precomH
h2 := <-wfd.precomH
o := <-wfd.precom0

// Step 2: (Commit)
//      Tell the 3 atoms to start bonding
h1 <- struct{}{}
h2 <- struct{}{}
o <- struct{}{}

// Step 3: (Postcommit)
//      Wait until the 3 atoms have finished before looping
// We re-use the same communication channel as (Commit)
<-h1
<-h2
<-o
```

```
type WaterFactoryWithDaemon struct {
    // Channels for atoms to send their arrival requests
    precomH chan chan struct{}
    precom0 chan chan struct{}
}
```

```
func (wfd *WaterFactoryWithDaemon) hydrogen(bond func()) {
    commit := make(chan struct{}) // Step 1: Create private communication channel
    wfd.precomH <- commit         // Step 2: (Precommit)
    <-commit                       // Step 3: (Commit)
    bond()                       // Step 4: Bond
    commit <- struct{}{}         // Step 5: (Postcommit)
}
```

Go Solution 1:

Issues:

- daemon thread is a single point of failure
- potential performance bottleneck due to scheduling decisions

Go Solution 1:

Issues:

- daemon thread is a single point of failure
- potential performance bottleneck due to scheduling decisions

Recall: 1 oxygen per H₂O, why not make oxygen the coordinator?

Go Solution 2:

```
func (wf *WaterFactoryWithLeader) oxygen(bond func()) {  
    // Step 1: Become leader  
    <-wf.oxygenMutex // For fun, we can use a channel as a mutex  
  
    // Step 2: (Precommit)  
    //     Receive arrival requests from 2 hydrogen atoms  
    h1 := <-wf.precomH  
    h2 := <-wf.precomH  
  
    // Step 3: (Commit)  
    //     Tell the 2 hydrogen atoms to start bonding  
    h1 <- struct{}{}  
    h2 <- struct{}{}  
  
    // Step 4: Bond  
    bond()  
  
    // Step 5: (Postcommit)  
    //     Wait until the 2 hydrogen atoms to finish  
    // We re-use the same communication channel as (Commit)  
    <-h1  
    <-h2  
  
    // Step 6: Step down from being leader  
    wf.oxygenMutex <- struct{}{}  
}
```

```
type WaterFactoryWithLeader struct {  
    oxygenMutex chan struct{}  
    precomH      chan chan struct{}  
}
```


```
func (wf *WaterFactoryWithLeader) hydrogen(bond func()) {  
    commit := make(chan struct{}) // Step 1: Create private communication channel  
    wf.precomH <- commit          // Step 2: (Precommit)  
    <-commit                       // Step 3: (Commit)  
    bond()                       // Step 4: Bond  
    commit <- struct{}{}         // Step 5: (Postcommit)  
}
```

Can still have issues if oxygen holding the “mutex” dies but the probability is lower

FIFO Semaphore


i.e. starvation-free

FIFO Semaphore(=2) Problem Statement




**E.g. Oxygen atoms
waiting to bond()**

FIFO Semaphore(=2) Problem Statement




**Semaphores don't
guarantee order!**

FIFO Semaphore(=2) Problem Statement




**Even if is sole consumer,
2 may be sleeping...**


FIFO Semaphore(=2) Problem Statement




FIFO Semaphore(=2) Problem Statement



FIFO Semaphore(=2) Problem Statement



FIFO Semaphore(=2) Problem Statement




FIFO using ticket queue

Demo 4


Demo 4: ticket queue

```
struct FIFOSemaphore {  
    std::atomic<std::ptrdiff_t> next_ticket{1};  
    std::atomic<std::ptrdiff_t> now_serving;  
  
    FIFOSemaphore(std::ptrdiff_t initial_count) : now_serving{initial_count}  
    {}  
  
    void acquire() {  
        auto my_ticket = next_ticket.fetch_add(1);  
        while (now_serving.load() < my_ticket) {}  
    }  
  
    void release() {  
        now_serving.fetch_add(1);  
    }  
};
```



Demo 4: ticket queue (initial_count = 2)




Demo 4: ticket queue (initial_count = 2)



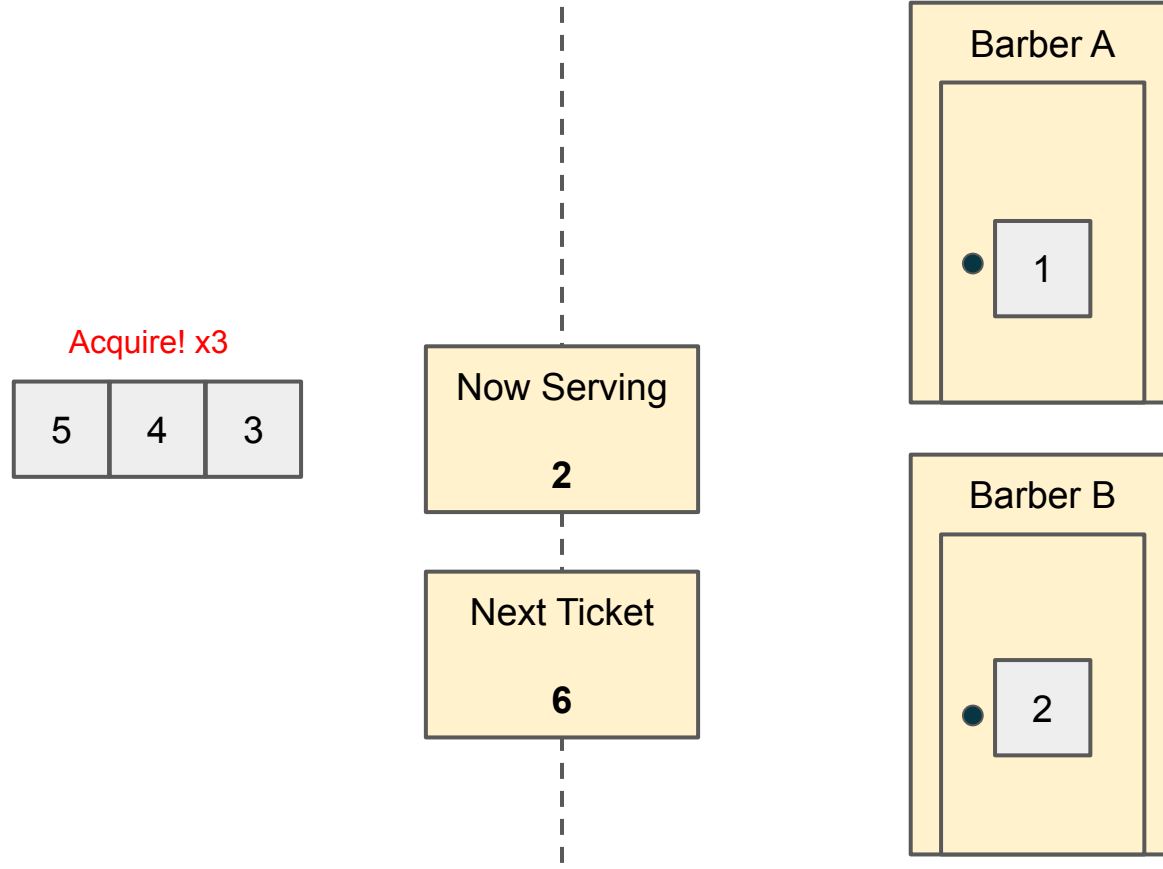
Demo 4: ticket queue (initial_count = 2)



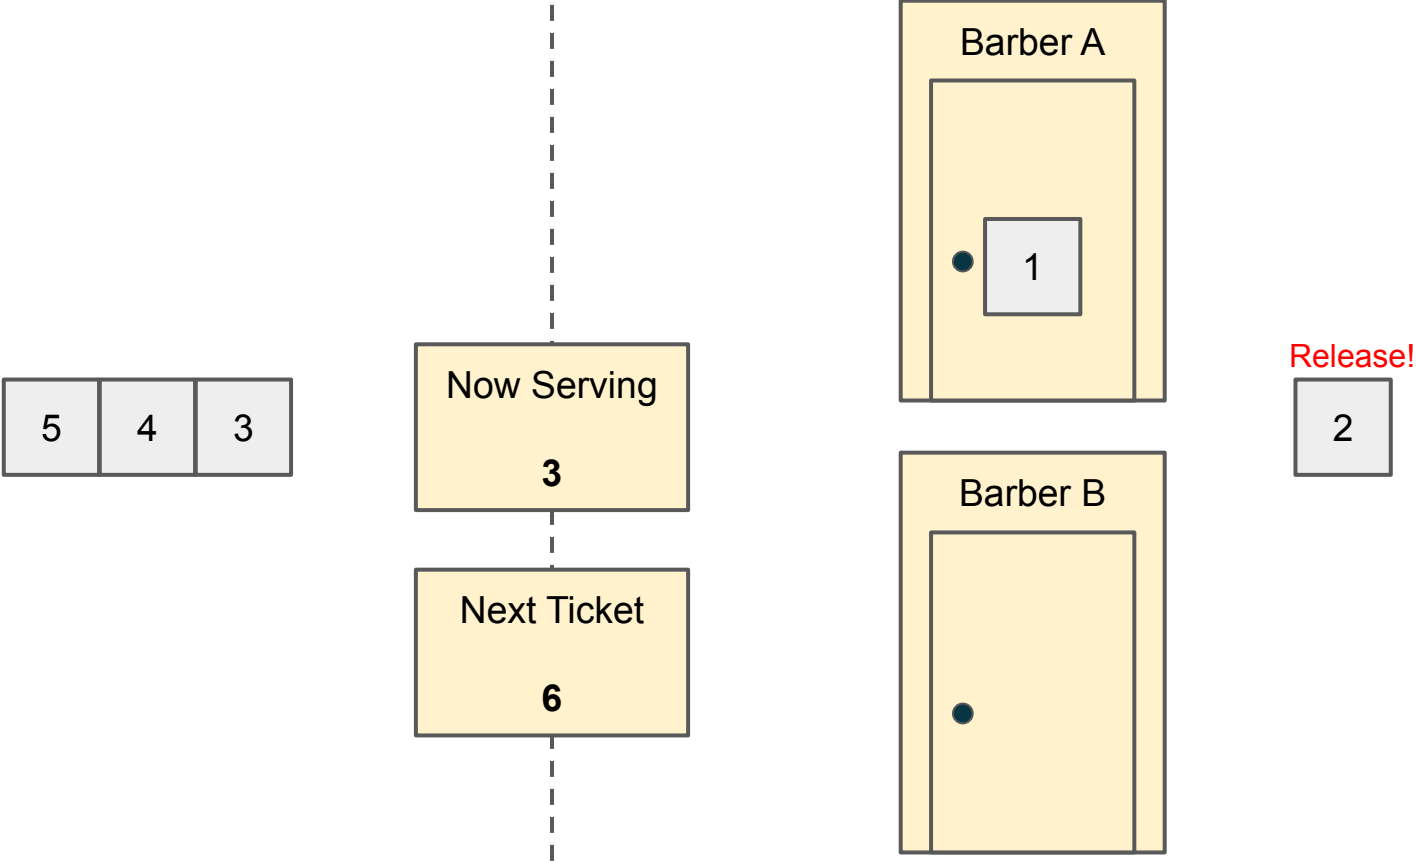
Demo 4: ticket queue (initial_count = 2)



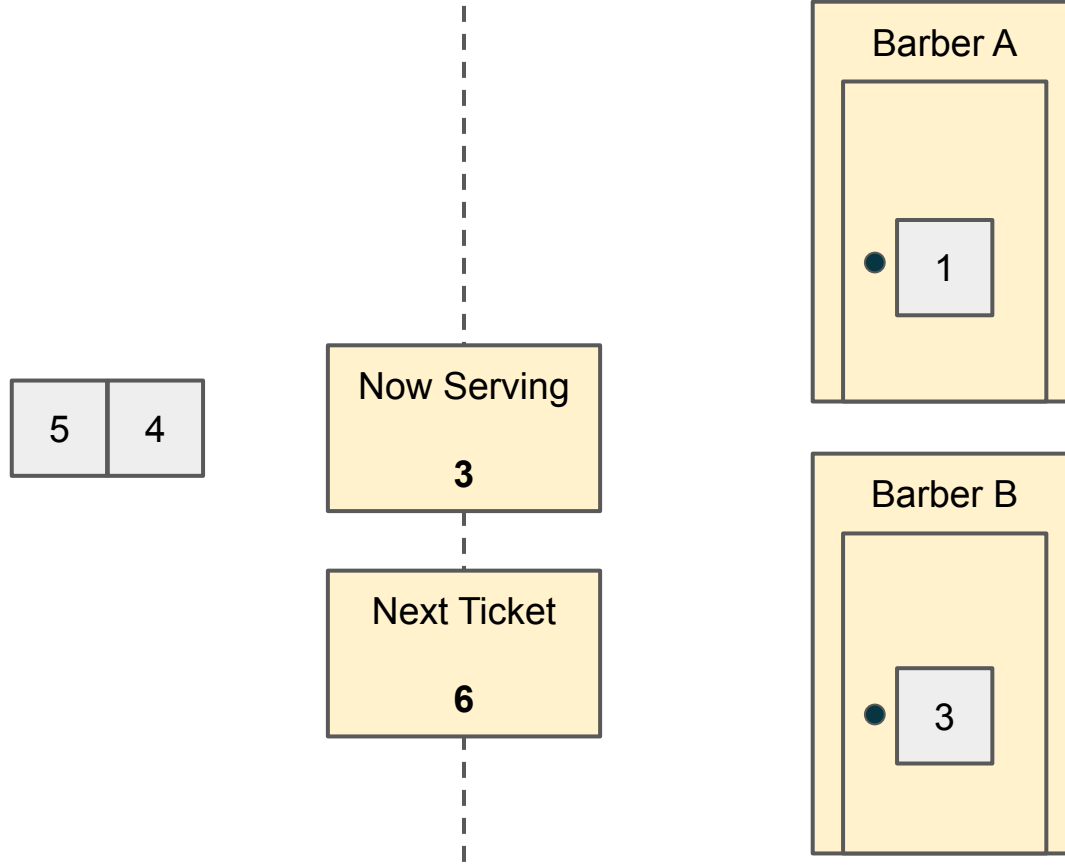
Demo 4: ticket queue (initial_count = 2)



Demo 4: ticket queue (initial_count = 2)



Demo 4: ticket queue (initial_count = 2)



Task 1: Replace spinwait

Try it yourself!

Use

condition_variable

or

atomic<T>::wait



```
struct FIFOSemaphore {
    std::atomic<std::ptrdiff_t> next_ticket{1};
    std::atomic<std::ptrdiff_t> now_serving;

    FIFOSemaphore(std::ptrdiff_t initial_count) : now_
    {}

    void acquire() {
        auto my_ticket = next_ticket.fetch_add(1);
        while(now_serving.load() < my_ticket) {}
    }

    void release() {
        now_serving.fetch_add(1);
    }
};
```


Task 1: Replace spinwait

```
struct FIFOSemaphore {  
    std::atomic<std::ptrdiff_t> next_ticket{1};  
    std::atomic<std::ptrdiff_t> now_serving;  
  
    FIFOSemaphore(std::ptrdiff_t initial_count) : now_serving{initial_count}  
    {}  
  
    void acquire() {  
        auto my_ticket = next_ticket.fetch_add(1);  
        auto old_now_serving = now_serving.load();  
        if (old_now_serving < my_ticket)   
            now_serving.wait(old_now_serving);  
    }  
  
    void release() {  
        now_serving.fetch_add(1);   
        now_serving.notify_one();  
    }  
};
```

Any problems?

Task 1: Replace spinwait

Performs atomic waiting operations. Behaves as if it repeatedly performs the following steps:

- Compare the value representation of `this->load(order)` with that of old.
 - If those are equal, then blocks until `*this` is notified by `notify_one()` or `notify_all()`, or the thread is unblocked spuriously.
 - Otherwise, returns.


These functions are guaranteed to return only if value has changed, even if underlying implementation unblocks spuriously.

```
if (old now_serving < my_ticket) {  
    now_serving.wait(old_now_serving, my_ticket);  
}  
  
void release() {  
    now_serving.fetch_add(1);  
    now_serving.notify_one();  
};
```

Just because old changed doesn't mean the condition is true





Task 1: Replace spinwait

`notify_one()` might not wake 3!



Task 1: Replace spinwait

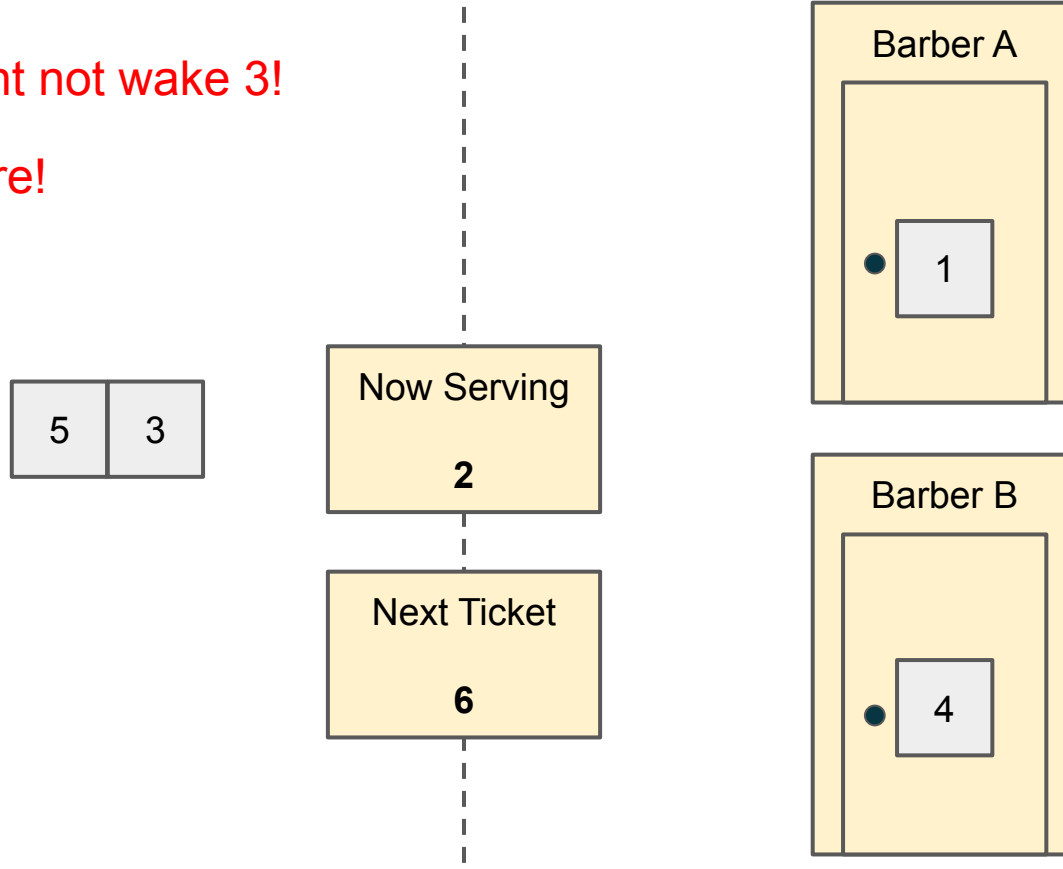
`notify_one()` might not wake 3!



Task 1: Replace spinwait

`notify_one()` might not wake 3!

-> not FIFO anymore!



Demo 5: Queue of semaphores

```
void acquire() {
    auto waiter = std::make_shared<Waiter>();
    {
        std::scoped_lock lock{mut};
        if (count > 0) {
            count--; // Positive count,
            return; // simply decrement without blocking
        }
        waiters.push(waiter); // Zero count, add to waiters
    }
    waiter->sem.acquire(); // and block on the semaphore
}
```

```
struct FIFOSemaphore5 {
    struct Waiter {
        std::binary_semaphore sem{0};
    };

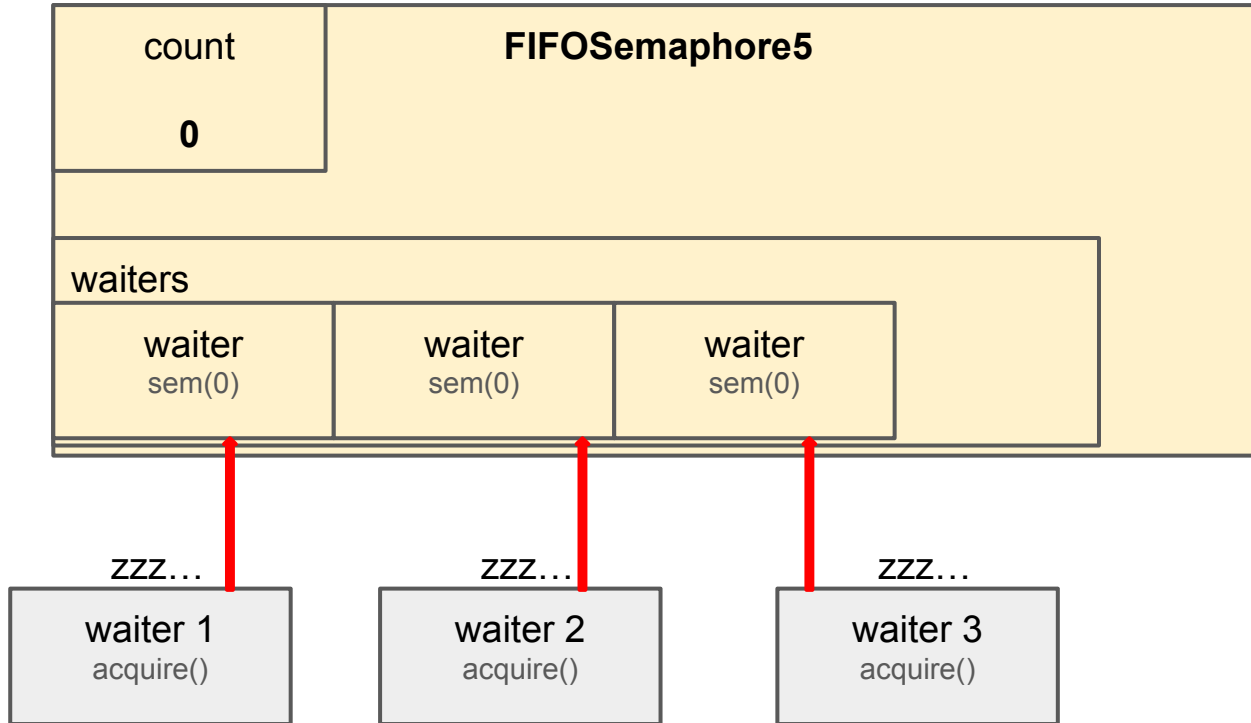
    std::mutex mut;
    std::queue<std::shared_ptr<Waiter>> waiters;
    std::ptrdiff_t count;

    FIFOSemaphore5(std::ptrdiff_t initial_count)
        : mut{}, waiters{}, count{initial_count} {}
}
```

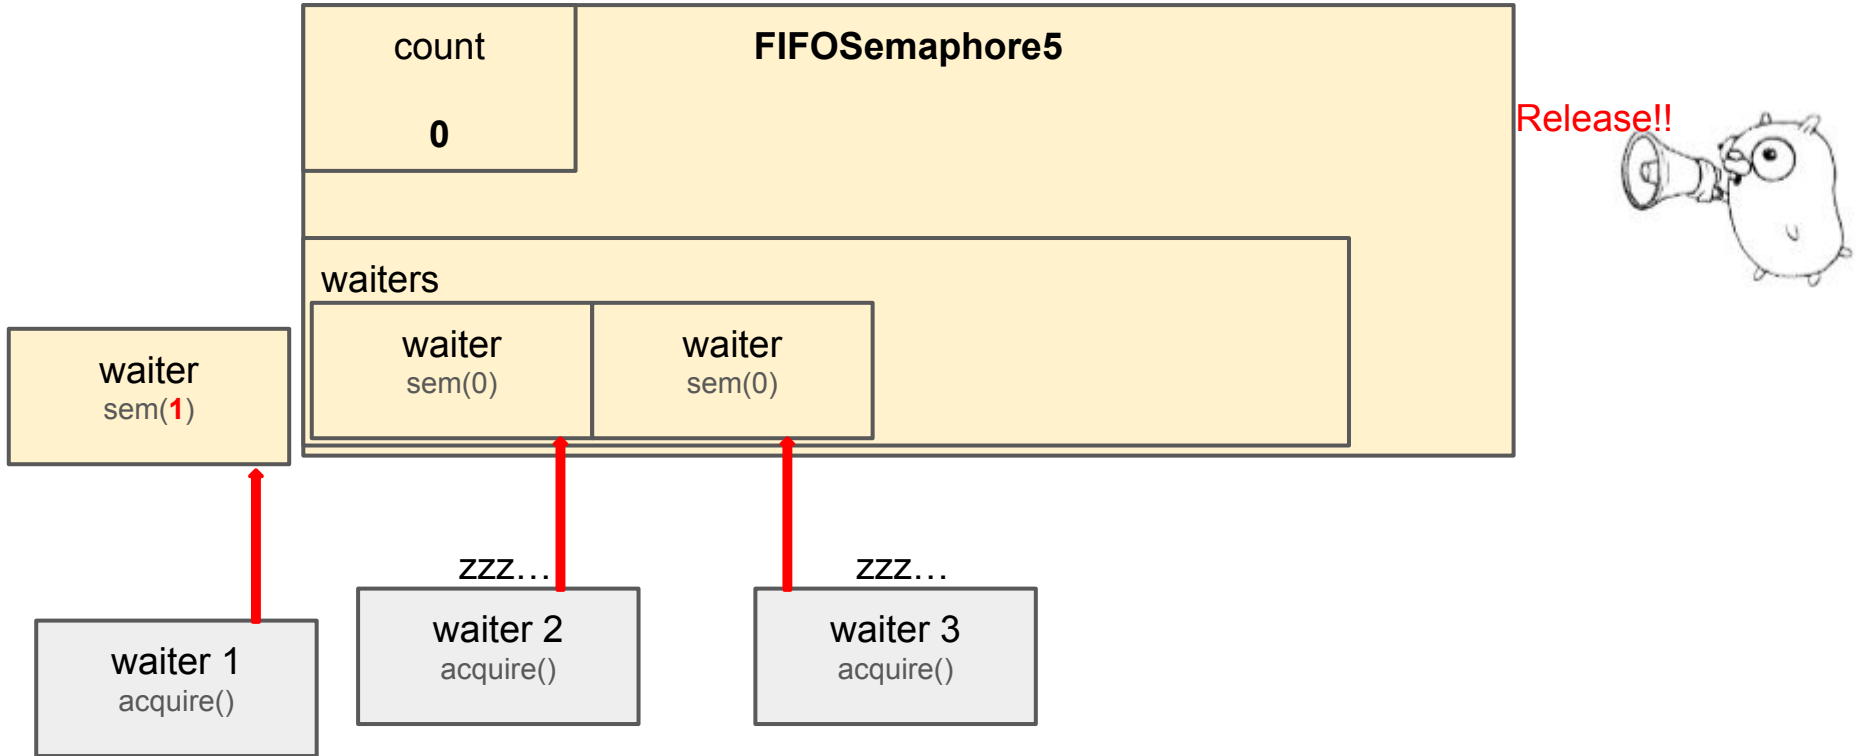
```
void release() {
    std::shared_ptr<Waiter> waiter;
    {
        std::scoped_lock lock{mut};
        if (waiters.empty()) {
            count++; // No waiters, simply increment count
            return;
        }

        waiter = waiters.front(); // Pop a waiter
        waiters.pop();
    }
    waiter->sem.release(); // and signal it
}
```

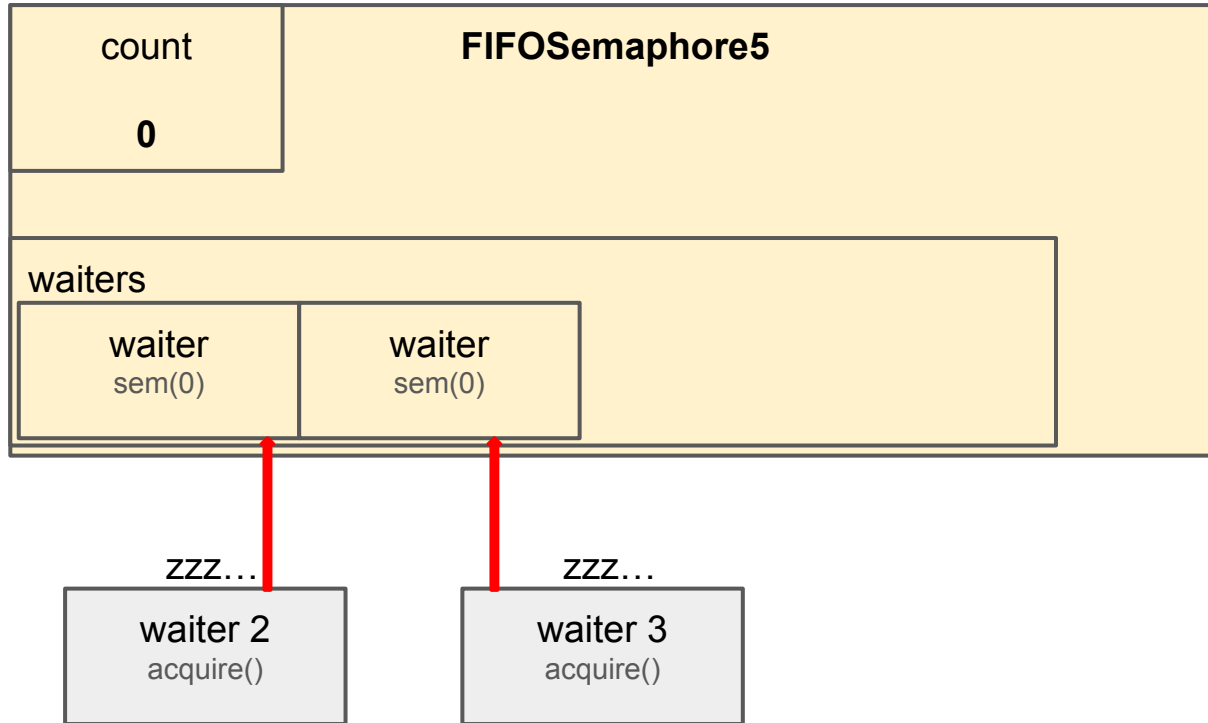
Demo 5: Queue of semaphores




Demo 5: Queue of semaphores



Demo 5: Queue of semaphores




Demo 5: Queue of semaphores



Release!!



Demo 5: Queue of semaphores



Demo 5: extras

Issues:

- 2 allocations each time waiter is added to queue

How about we do only 1 allocation?

Demo 5: extras

```
void acquire() {
    Waiter waiter;

    {
        std::scoped_lock lock{mut};
        if (count > 0) {
            count--;
            return;
        }

        if (back == nullptr) {
            front = back = &waiter;
        } else {
            back->next = &waiter;
            back = &waiter;
        }
    }
    waiter.sem.acquire();
}
```

Since destruction of waiter is after all “external” accesses to it, no invalid memory access.
tldr; no lifetime issues

```
struct FIFOSemaphore7 {
    struct Waiter {
        std::binary_semaphore sem{0};
        Waiter *next{nullptr};
    };

    std::mutex mut;
    Waiter *front;
    Waiter *back;
    std::ptrdiff_t count;

    FIFOSemaphore7(std::ptrdiff_t initial_count)
        : mut{}, front{nullptr}, back{nullptr}, count{initial_count} {}

    void release() {
        Waiter *head;
        {
            std::scoped_lock lock{mut};
            if (front == nullptr) {
                count++;
                return;
            }
            head = front;
            front = front->next;
            if (front == nullptr) {
                back = nullptr;
            }
        }
        head->sem.release();
    }
}
```

Demo 6: Using a buffered channel

- Sending on a full/unbuffered channel -> block
- Suppose `initial_count = N`
 - Make a buffered channel of size N.
 - If $M > N$ goroutines arrive at the same time, $(M - N)$ goroutines block.
 - release -> send to channel
 - acquire -> recv from channel

Basically a semaphore!

Demo 6: Using a buffered channel

```
type hchan struct {
    qcount    uint           // total data in the queue
    dataqsiz  uint           // size of the circular queue
    buf       unsafe.Pointer // points to an array of dataqsiz elements
    elemsize  uint16
    closed    uint32
    elemtype  *_type // element type
    sendx     uint     // send index
    recvx     uint     // receive index
    recvq     waitq    // list of recv waiters
    sendq     waitq    // list of send waiters
}

type waitq struct {
    first *sudog
    last  *sudog
}
```

```
type sudog struct {
    // The following fields are protected by the hchan.lock of the
    // channel this sudog is blocking on. shrinkstack depends on
    // this for sudogs involved in channel ops.

    g *g

    next *sudog
    prev *sudog
}
```

current impl is fifo, but not guaranteed by the standard.

Demo 7: Using a daemon goroutine

Bad idea to rely on implementation-defined behaviour.

How about we have a goroutine manage others!

Demo 7: Using a daemon goroutine

```
type Semaphore2 struct {  
    acquireCh chan chan struct{}  
    releaseCh chan struct{}  
}
```

```
func (s *Semaphore2) Acquire() {  
    ch := make(chan struct{})  
    // Send daemon a channel that can be used to unblock us  
    s.acquireCh <- ch  
    // Block until daemon decides to unblock us  
    <-ch  
}
```

```
func (s *Semaphore2) Release() {  
    s.releaseCh <- struct{}{}  
}
```

```
func NewSemaphore2(initial_count int) *Semaphore2 {  
    sem := new(Semaphore2)  
    sem.acquireCh = make(chan chan struct{}, 100)  
    sem.releaseCh = make(chan struct{}, 100)  
  
    go func() {  
        count := initial_count  
        // The FIFO queue that stores the channels used to unblock waiters  
        waiters := NewChanQueue()
```

Demo 7: Using a daemon goroutine

```
for {  
    select {  
    case <-sem.releaseCh:  
        // ...  
  
    case ch := <-sem.acquireCh:  
        // ...  
    }  
}
```

Either process a release or acquire request

Demo 7: Using a daemon goroutine

```
case ch := <-sem.acquireCh: // Decrement or add a waiter
    if count > 0 {
        count--
        ch <- struct{}{} // Since count is +ve, don't block the waiter
    } else {
        waiters.PushBack(ch) // Add waiter to back of the wait queue
    }
}
```

Either process a release or acquire request

Demo 7: Using a daemon goroutine

```
case <-sem.releaseCh: // Increment or unblock a waiter
    if waiters.Len() > 0 {
        ch := waiters.Pop()
        ch <- struct{}{} // Unblocks the oldest waiter
    } else {
        count++
    }
```

Either process a release or acquire request

Demo 7: Using a daemon goroutine

Why a buffer of 100?

```
func NewSemaphore2(initial_count int) *Semaphore2 {  
    sem := new(Semaphore2)  
    sem.acquireCh = make(chan chan struct{}, 100)  
    sem.releaseCh = make(chan struct{}, 100)  
  
    go func() {  
        count := initial_count  
        // The FIFO queue that stores the channels used to unblock waiters  
        waiters := NewChanQueue()
```

Demo 7: Using a daemon goroutine

Why a buffer of 100?

If goroutine blocks on

```
func (s *Semaphore2) Acquire() {  
    ch := make(chan struct{})  
    // Send daemon a channel that can be used to  
    unblock us  
    s.acquireCh <- ch  
    // Block until daemon decides to unblock us  
    <-ch  
}
```

for a long time (wakeups not in FIFO), no guarantees of FIFO.

BUT, if send succeeds (buffered in channel), we get FIFO behavior.

```
type hchan struct {  
    qcount    uint           // total data in the queue  
    dataqsiz  uint           // size of the circular queue  
    buf       unsafe.Pointer // points to an array of dataqsiz elements  
    elemsize  uint16  
    closed    uint32  
    elemtype  *_type // element type  
    sendx     uint       // send index  
    recvx     uint       // receive index  
    recvq     waitq      // list of recv waiters  
    sendq     waitq      // list of send waiters  
  
    // lock protects all fields in hchan, as well as several  
    // fields in sudogs blocked on this channel.  
    //  
    // Do not change another G's status while holding this lock  
    // (in particular, do not ready a G), as this can deadlock  
    // with stack shrinking.  
    lock mutex  
}
```

Pls scan for attendance



See you next week!